

# Processes

Mythili Vutukuru

CSE, IIT Bombay

# Recap: OS manages hardware for users

- Creates and runs processes on a system
- Allocates memory to processes, manages virtual address space
- Schedules processes on CPU, and switches between processes to timeshare CPU
- Handles interrupts from I/O devices, and other events

# The process abstraction

- Process is a running program
- When program is run, OS creates new process, allocates memory, initializes CPU context, and starts process in user mode
- User program runs on CPU normally, unless OS needs to step in for system calls, interrupts, ...

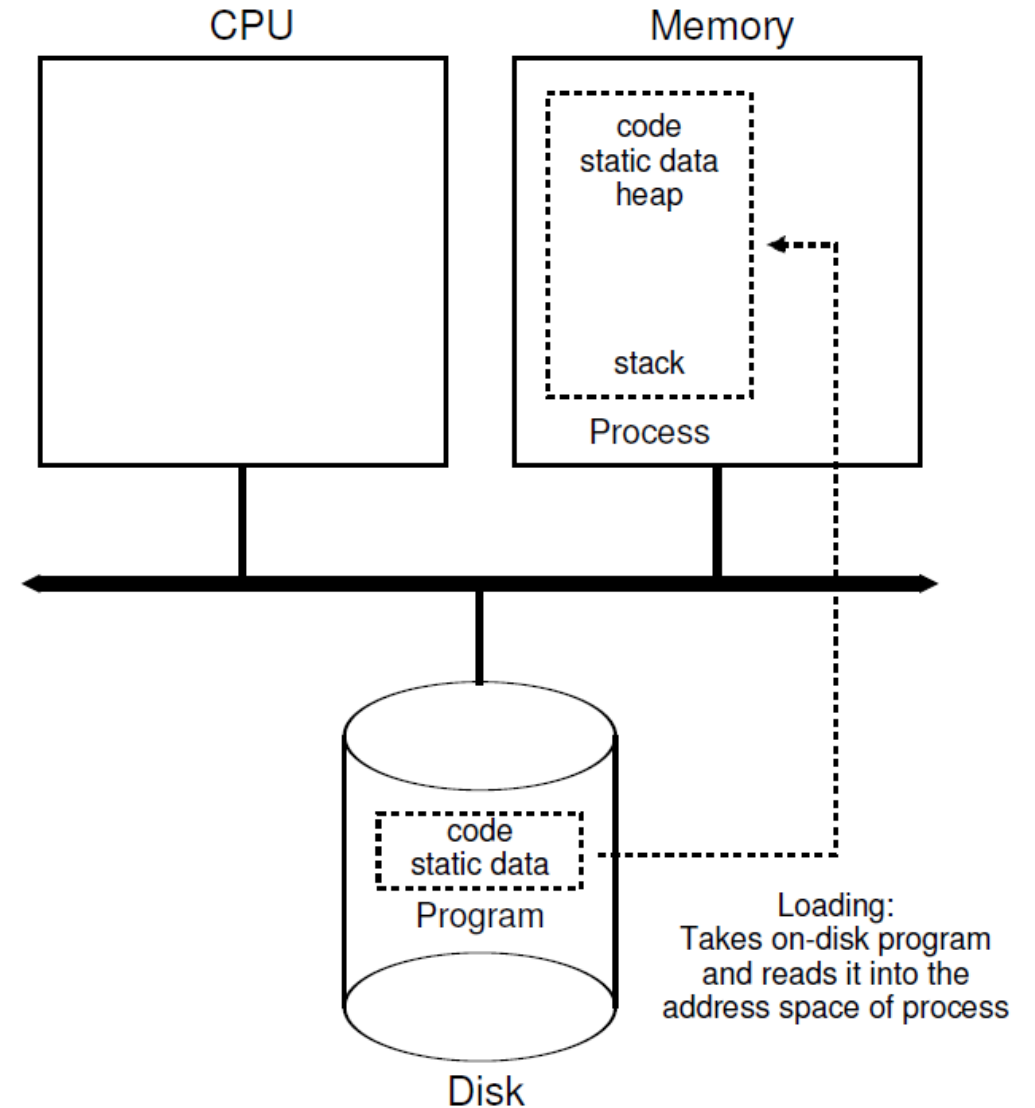


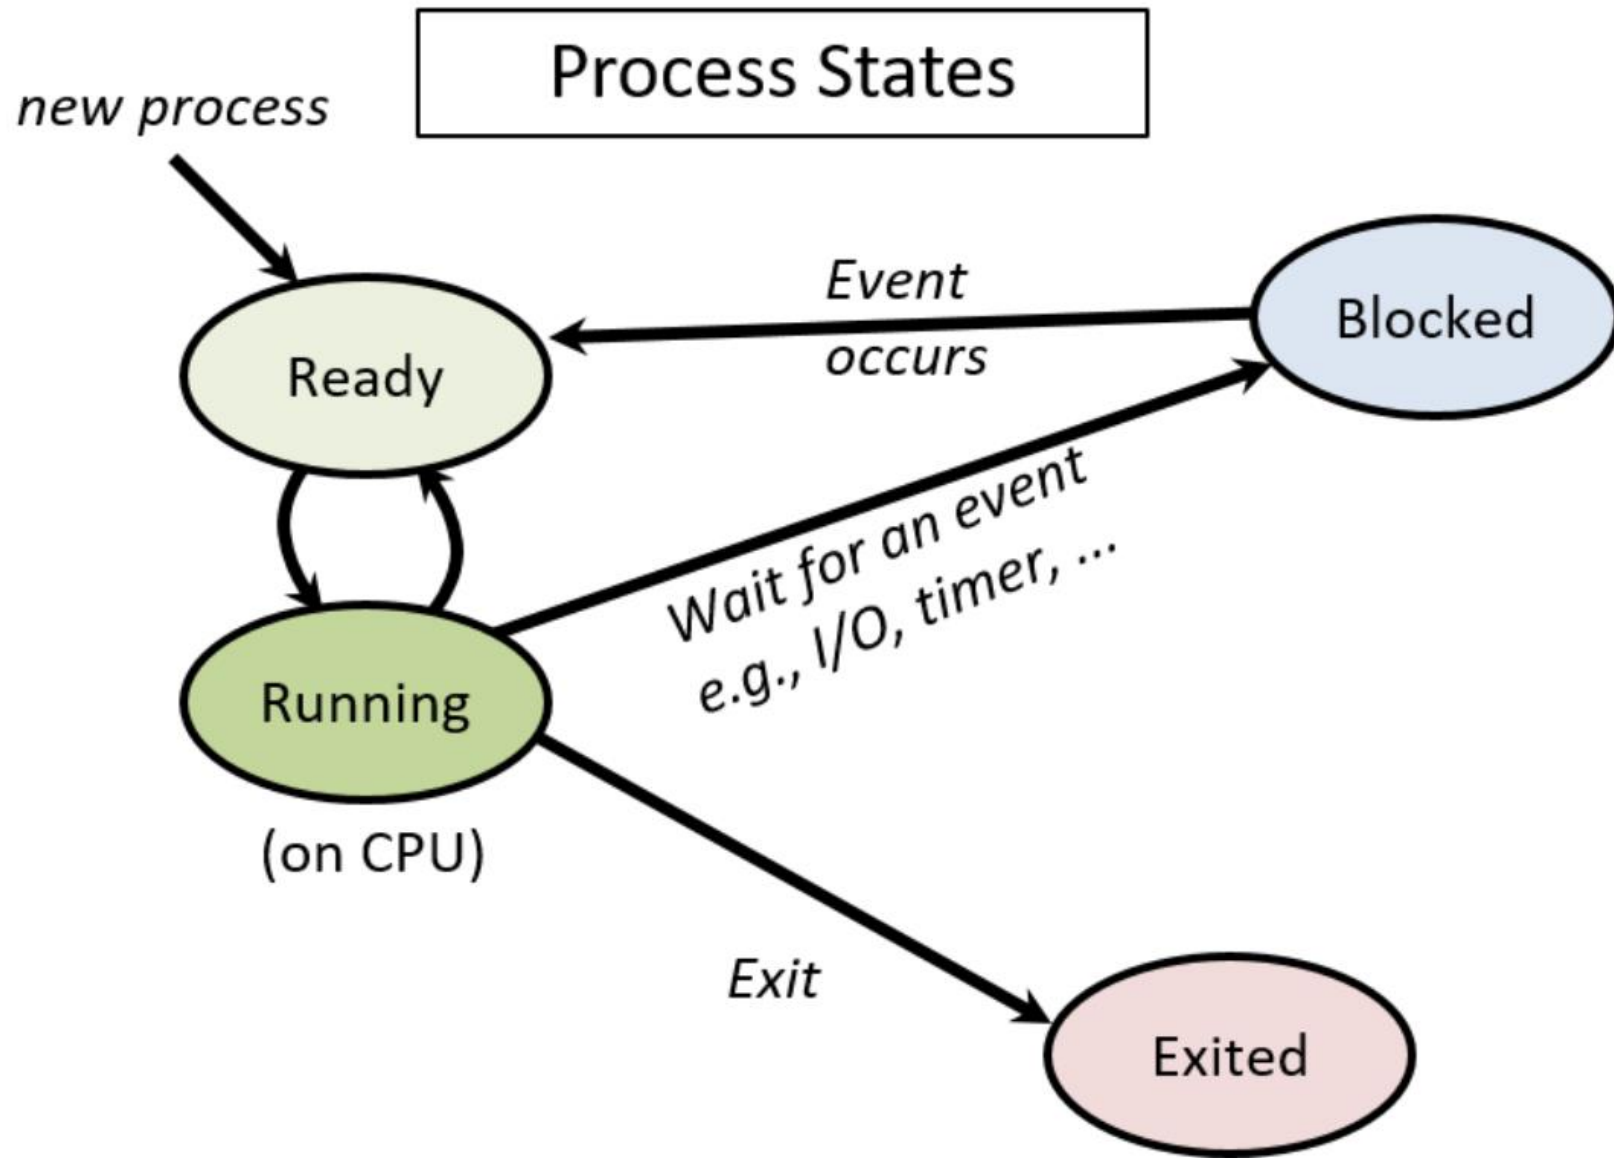
Figure 4.1: Loading: From Program To Process

# What defines a process?

- Every process has a unique **process identifier (PID)**
- Process occupies some memory in RAM (**memory image**)
  - Code+data from executable
  - Stack+heap allocated for runtime use
- The execution **context** of the process (values of CPU registers)
  - PC has address of instruction of process, some registers have process data
  - Process context is in CPU registers when process is running on CPU
  - Context saved in memory when process is paused, restored when run again
- Ongoing **communication with I/O devices**
  - Information is maintained about files that are open, ongoing network connections, other active connections to I/O devices

# States of a process

- OS manages multiple active processes at the same time. An active process can be in one of the following situations.
- **Running**: currently executing on CPU
  - CPU registers contain context of process
- **Blocked/suspended/sleeping**: process cannot run for some time
  - Example: process has requested data from disk, command issued, but process cannot proceed until the data from disk is available
- **Ready/runnable**: ready to run but waiting for OS scheduler to switch the process in
  - Many processes can be ready but scheduler can only run one on a CPU core
- Context of blocked and ready processes is saved in memory, so that they can continue to run later on



*Figure 1. The states of a process during its lifetime*

# Example: process state transitions

- Consider a system that has two user processes P0 and P1
  - Initially P0 is running, P1 is ready and awaiting its turn
  - P0 wants to read a file from disk via a system call
  - OS handles the system call and gives command to disk, but data is not available immediately
  - Process P0 is moved to blocked state, OS switches to process P1
  - Process P1 runs for some time, and then an interrupt occurs from disk
  - CPU jumps to OS which handles interrupt, P0 is moved to ready state
  - OS can continue to run P1 again after interrupt and OS scheduler switches to ready process P0 later on after some time

# Example: process state transitions

Time	Process <sub>0</sub>	Process <sub>1</sub>	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process <sub>0</sub> initiates I/O
4	Blocked	Running	Process <sub>0</sub> is blocked,
5	Blocked	Running	so Process <sub>1</sub> runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process <sub>1</sub> now done
9	Running	–	
10	Running	–	Process <sub>0</sub> now done

Figure 4.4: Tracing Process State: CPU and I/O



# Process State Transitions

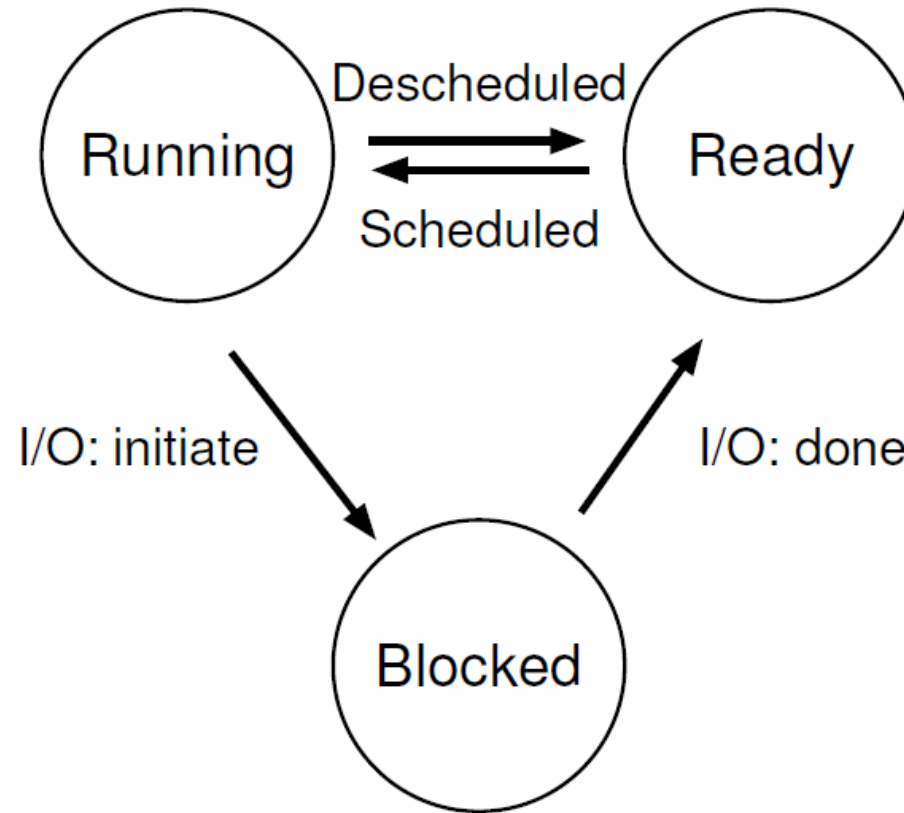


Figure 4.2: **Process: State Transitions**

# Process control block (PCB)

- All information about a process is stored in a kernel data structure called the **process control block (PCB)**
  - Process identifier (PID)
  - Process state (running, ready, blocked, terminated, ..)
  - Pointers to other related processes (parent, children)
  - Saved CPU context of process when it is not running
  - Information related to memory locations of a process
  - Information related to ongoing I/O communication
  - ...
- PCB is known by different names in different OS
  - struct proc in xv6
  - task\_struct in Linux

# PCB in xv6: struct proc

```
2334 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
2335
2336 // Per-process state
2337 struct proc {
2338     uint sz;                // Size of process memory (bytes)
2339     pde_t* pgdir;           // Page table
2340     char *kstack;           // Bottom of kernel stack for this process
2341     enum procstate state;    // Process state
2342     int pid;                // Process ID
2343     struct proc *parent;     // Parent process
2344     struct trapframe *tf;    // Trap frame for current syscall
2345     struct context *context; // swtch() here to run process
2346     void *chan;              // If non-zero, sleeping on chan
2347     int killed;              // If non-zero, have been killed
2348     struct file *ofile[NOFILE]; // Open files
2349     struct inode *cwd;       // Current directory
2350     char name[16];           // Process name (debugging)
2351 };
2352
```

# struct proc: page table

- Every instruction or data item in the memory image of process (code/data, stack, heap, etc.) has an address
  - Virtual addresses, starting from 0
  - Actual physical addresses in memory can be different (all processes cannot store their first instruction at address 0)
- Page table of a process maintains a mapping between the virtual addresses and physical addresses
- Page table used to translate virtual addresses to physical addresses

# struct proc: kernel stack

- Stack to store CPU context when process jumps to kernel mode from user mode, or when process is context switched out
  - Why separate stack? OS does not trust user stack
  - Separate area of memory in the kernel, not accessible by regular user code
  - Linked from struct proc of a process

# struct proc: list of open files

- Array of pointers to open files
  - When user opens a file, a new entry is created in this array, and the index of that entry is passed as a file descriptor to user
  - Subsequent read/write calls on a file use this file descriptor to refer to the file
  - First 3 files (array indices 0,1,2) open by default for every process: standard input, output and error
  - Subsequent files opened by a process will occupy later entries in the array

# Process table (ptable) in xv6

- Ptable in xv6 is a fixed-size array of all processes
- Real kernels have dynamic-sized data structures

```
2409 struct {  
2410     struct spinlock lock;  
2411     struct proc proc[NPROC];  
2412 } ptable;
```

# CPU scheduler in xv6

- The OS loops over all runnable processes in ptable, picks one, and sets it running on the CPU

```
2768     // Loop over process table looking for process to run.
2769     acquire(&ptable.lock);
2770     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2771         if(p->state != RUNNABLE)
2772             continue;
2773
2774         // Switch to chosen process.  It is the process's job
2775         // to release ptable.lock and then reacquire it
2776         // before jumping back to us.
2777         c->proc = p;
2778         switchvm(p);
2779         p->state = RUNNING;
```



# Process state transition: example in xv6

- A process that needs to sleep (e.g., for disk I/O) will set its state to SLEEPING and invoke scheduler
- Scheduler will run its loop and find another ready process to run

```
2873 void
2874 sleep(void *chan, struct spinlock *lk)
2875 {
2876     struct proc *p = myproc();
2877
2878     if(p == 0)
2879         panic("sleep");
2880
2881     if(lk == 0)
2882         panic("sleep without lk");
2883
2884     // Must acquire ptable.lock in order to
2885     // change p->state and then call sched.
2886     // Once we hold ptable.lock, we can be
2887     // guaranteed that we won't miss any wakeup
2888     // (wakeup runs with ptable.lock locked),
2889     // so it's okay to release lk.
2890     if(lk != &ptable.lock){
2891         acquire(&ptable.lock);
2892         release(lk);
2893     }
2894     // Go to sleep.
2895     p->chan = chan;
2896     p->state = SLEEPING;
2897
2898     sched();
2899 }
```

# Booting

- What happens when you boot up a computer system?
- Basic Input Output System (BIOS) starts to run
  - Resides in non-volatile memory, sets up all other hardware
- BIOS locates the boot loader in the boot disk (hard disk, USB, ..)
  - Simple program whose job is to locate and load the OS
  - Present in the first sector of the boot disk
  - Combination of assembly and C code
- Boot loader sets up CPU registers suitably, loads kernel image from disk to memory, transfers control to kernel
- OS code starts to run, exposes terminal to user, user starts programs

# Booting real systems

- Bootloader must fit into 512 bytes (first sector of boot disk) to be found easily by BIOS
- Bootloaders for simple/old OS could fit into one sector, but no longer the case for modern OS
- Real life bootloaders are complex, need to read a large kernel image from disk and network, do not fit into 512 bytes
- Real life booting is two step process: BIOS loads simple bootloader, which loads a more complex bootloader, which then loads the OS

# Summary

*Three ready processes*

*Three ready processes*

*See how they run*

*See how they run*

*They all sat in the scheduler's queue*

*The scheduler sent them out to execute*

*They made an I/O request and went to sleep*

*Became ready to run after interrupts beeped*

*Did you ever see such a sight in your life*

*As three ready processes*

(To be sung in the tune of the nursery rhyme “Three Blind Mice”)