

Inter-process communication (IPC)

Mythili Vutukuru
CSE, IIT Bombay

Why inter-process communication

- Application logic in a single system is often distributed across multiple processes: why?
 - Different processes developed independently by different teams
 - Different programming languages and frameworks used for different tasks
- Processes in a system do not share any memory with each other by default, so how do they communicate information with each other?
 - Cannot share variables or data structures in programs across processes
 - Parent and child have identical but separate memory images after fork, changes made in one process and not seen by other
- **Inter-process communication (IPC)** mechanisms, available via operating system syscalls, allow processes to exchange information

Example: web application architecture

- Example: web applications typically composed of multiple processes
- **Web server** process handles HTTP (web) requests/responses
 - Written in a language like C/C++ for high performance
 - Returns responses for static content directly by reading files from disk
- Requests needing dynamic response are handled by **application server**
 - App server parses HTTP requests, generates HTTP response according to the business logic specified by user, sends response back to client via web server
 - Scripting languages may be used for easy text parsing and manipulation
- Application server stores/retrieves app data in a **database**
- Several web application frameworks available to build web applications having such architectures, e.g., Python Django, React etc.

IPC mechanisms

- **Unix domain sockets**: processes open sockets, send and receive messages to each other via socket system calls
- **Message queues**: sender posts a message to a mailbox, receiver retrieves message later on from mailbox
- **Pipes**: unidirectional communication channel between two processes
- **Shared memory**: same physical memory frame mapped into virtual address space of multiple processes in order to share memory
- Different IPC mechanisms are useful in different scenarios

Sockets

- Sockets = abstraction to communicate between two processes
 - Each process opens socket, and pair of sockets can be connected
 - Client-server paradigm: one process opens socket first (**server**) and another process connects its socket to the first one (**client**)
 - One process writes a message into one socket, another process can read it, and vice versa (**bidirectional communication**)
 - Processes can be in same machine or on different machines
 - If processes on same machine, messages stored temporarily in OS memory before delivering to destination process
 - If processes on different machines, messages sent over network

Types of sockets (1)

- **Unix domain (local) sockets** are used to communicate between processes on the same machine
- **Internet sockets** are used to communicate between processes in different machines
- Local sockets identified by a pathname, Internet sockets identified by IP (Internet Protocol) address and port number
- Client and server sockets differentiated by who starts first and who connects later: server sockets started first on a well-known “address”, client process connects to server using the server address

Types of sockets (2)

- **Connection-based sockets**: one client socket and one server socket are explicitly connected to each other
 - After connection, the two sockets can only send and receive messages to each other
- **Connection-less sockets**: one socket can send/receive messages to/from multiple other sockets
 - Address of other endpoint can be mentioned on each message
- Type of socket (local or internet, connection-oriented or connection-less) is specified as arguments to system call that creates sockets

Creating a socket

```
sockfd = socket(...)  
bind(sockfd, address)
```

- System call “socket” used to create a socket
 - Takes type of socket as arguments
 - Returns **socket file descriptor** (similar to file descriptor when file is opened)
 - Used as handle for all future operations on the socket
- A socket can optionally **bind to an address** (pathname or IP address/port number) using “bind” system call
 - Server sockets bind to well known address, so that clients can connect
 - Client sockets need not bind, OS can assign temporary address
- Close system call closes a socket when done

Data exchange using connection-less sockets

- Function `sendto` is used to send a message from one socket to another connection-less socket in another process
 - Arguments: socket fd, message to send, address of remote socket
- Function `recvfrom` is used to receive a message from a socket
 - Arguments: socket fd, message buffer into which received message is copied, socket address structure into which address of remote endpoint is filled
 - When a process receives a message on connection-less socket, it can find out the address of other endpoint, and use this address to reply back

Client

```
sockfd = socket(..)  
char message[1024]  
sendto(sockfd, message, server_sockaddr, ..)
```

Server

```
sockfd = socket(..)  
bind(sockfd, server_address)  
recvfrom(sockfd, message, client_sockaddr, ..)
```

Connecting sockets

- Connection-oriented sockets must be explicitly connected to each other before exchanging messages
- After server binds socket to well-known address, it uses “listen” system call to make the socket **listen** for new connections
- Client uses “connect” system call to **connect** to a server listen socket
- Server uses “accept” system call to **accept** new connection requests
 - Returns a new socket file descriptor to communicate exclusively with a client
- At server: one **listen socket** to accept new connections, one **connected socket** for every connected client to send/recv messages

Client

```
sockfd = socket(..)  
connect(sockfd, server_sockaddr, ..)
```

Server

```
sockfd = socket(..)  
bind(sockfd, server_address)  
listen(sockfd, ..)  
newsockfd = accept(sockfd, ..)
```

Data exchange using connected sockets

- After client connects to server, pair of sockets used to exchange data
 - Note that per-client connected socket is used at server, not listen socket
 - System calls send/write used to send message on a connected socket
 - System calls recv/read used to receive message on a connected socket
- Arguments to send/recv: socket fd, message buffer, buffer length, flags
 - Return value is number of bytes read/written or error
 - No need to specify socket address on every message, as connected already

Client

```
sockfd = socket(..)
connect(sockfd, server_sockaddr, ..)
n = send(sockfd, req_buf, req_len, ..)
n = recv(sockfd, resp_buf, resp_len, ..)
```

Server

```
sockfd = socket(..)
bind(sockfd, server_address)
listen(sockfd, ..)
newsockfd = accept(sockfd, ..)
n = recv(newsockfd, req_buf, req_len, ..)
n = send(newsockfd, resp_buf, resp_len, ..)
```

Message queues

```
msgid = msgget(key, ...)  
msgsnd(msgid, message, ...)  
msgrcv(msgid, message, ...)
```

- Message queues used for exchanging messages between processes
 - Open connection to message queue identified by a “key”, get a handle
 - Sender opens connection to message queue, sends message
 - Receiver opens connection to message queue, retrieves message later on
 - Message buffered within message queue / mailbox until retrieved by receiver
- Example: IPC in web application using message queues
 - Web server posts dynamic HTTP requests into message queue
 - App server retrieves requests and processes them
 - App server posts responses into message queue for web server

Pipes

```
int fd[2]
pipe(fd) //anonymous
read(fd[0], message, ..)
write(fd[1], message, ..)
```

```
mkfifo(name, ..)
fd0 = open(name, O_RDONLY)
read(fd0, message, ..)
fd1 = open(name, O_WRONLY)
write(fd1, message, ...)
```

- Pipe is a **unidirectional FIFO channel** into which bytes are written at one end, read from other end
 - System call “pipe” creates a pipe channel, with two file descriptors for endpoints
 - One file descriptor used to write into pipe, one to read from pipe
 - Data written into pipe is stored in a buffer of the pipe channel until read
 - Bi-directional communication needs two pipes
- **Anonymous pipes** only available for use within process and its children
 - Pipe file descriptors point to same pipe structure in parent and child
- How to use pipes between unrelated processes? **Named pipes**
 - Named pipes opened with a pathname, accessible across processes
 - One process accesses read end of pipe, another opens write end

Blocking vs. non-blocking IPC

- Same high level concept across sockets, pipes, message queues
 - Sender sends message, temporarily stored in some memory inside OS
 - Receiver retrieves message later on from temporary OS memory
- Send/receive system calls can **block**
 - Sender can block if temporary buffer is full
 - Receiver can block if temporary buffer is empty
- Possible to configure IPC to be **non-blocking** using syscalls
 - Send/receive will return with error instead of blocking

Shared memory

```
shmid = shmget(key, ..)  
char *data = shmat(shmid, ..)
```

- Processes in a system do not share any memory by default
 - Child process gets copy of parent memory image, modifies independently
- Shared memory: a way for two processes to share memory
 - Same memory appears in memory image of multiple processes
 - Shared memory **segment** identified by a **unique key**
 - Process can request to map or “attach” a specific shared memory segment into its memory image by using key
- Processes may need extra mechanisms for coordination besides shared memory
 - E.g., how does one process know when another process has modified shared memory?

Summary

*One process P1 said to its friend P2
"It's been long since I have heard from you
Let's communicate by sharing some memory?
Or how about a socket, a pipe, or a message queue?"*

*"No", said P2, "none of this is easy
I do not know your port number or unique key
Neither are we related to use the same pipe
So I really do not know how you can reach me"*

*Processes, like humans, always have an excuse
To avoid talking, we are never short of a ruse
But find a way out, get the conversation going
Nothing's better than IPC* to get over the blues!*

* IPC = Inter Person Communication