

Memory management system call implementation in xv6

Mythili Vutukuru
CSE, IIT Bombay

Memory allocation to user processes in xv6

- How does OS implement various memory system calls?
- Understand how OS allocates memory to user processes using xv6 as example (real life OS have a much more complicated story)
- xv6 OS maintains list of free pages available in DRAM
 - All memory in $[0, \text{PHYSTOP}]$ not used by OS is added to free list
- Whenever new process memory image (address space) needs to be created during system calls (e.g., fork, exec), OS allocates pages from free list to process, updates its page table
- Processes can also request more pages from OS using sbrk system call
 - sbrk invoked by malloc to expand heap
 - xv6 has no mmap system call

Memory allocation in the kernel

- OS needs memory for its data structures, must allocate it from its free pages only
- For large allocations, OS allocates a page for itself
- For smaller allocations, OS implements its own versions of slab allocator or buddy allocator
 - Cannot use libc and malloc in kernel!
 - Slab allocator for common data structures, e.g., slab of PCBs
 - Buddy allocator for variable sized allocations
 - OS does not use very general variable sized allocation for efficiency reasons
- xv6 uses only page sized allocations, other data structures are fixed size. e.g., ptable

Maintaining free memory in xv6

- After boot up, RAM contains OS code/data and free pages (physical memory frames)
- OS collects all free pages into a free list, so that they can be allocated to user processes
- Free list is a linked list, pointer to next free page embedded within previous free page
- Kernel maintains pointer to first page in the free list
- Pages from free list allocated for code/data/stack/heap as well as page table of process

```
3115 struct run {  
3116     struct run *next;  
3117 };  
3118  
3119 struct {  
3120     struct spinlock lock;  
3121     int use_lock;  
3122     struct run *freelist;  
3123 } kmem;
```

Managing free pages in xv6: kalloc and kfree

- Anyone who needs a free page calls kalloc()
 - Sets free list pointer to next page and returns first free page on list
- When memory needs to be freed up, kfree() is called
 - Add free page to head of free list, update free list pointer

```
3186 char*
3187 kalloc(void)
3188 {
3189     struct run *r;
3190
3191     if(kmem.use_lock)
3192         acquire(&kmem.lock);
3193     r = kmem.freelist;
3194     if(r)
3195         kmem.freelist = r->next;
3196     if(kmem.use_lock)
3197         release(&kmem.lock);
3198     return (char*)r;
3199 }
```

```
3163 void
3164 kfree(char *v)
3165 {
3166     struct run *r;
3167
3168     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
3169         panic("kfree");
3170
3171     // Fill with junk to catch dangling refs.
3172     memset(v, 1, PGSIZE);
3173
3174     if(kmem.use_lock)
3175         acquire(&kmem.lock);
3176     r = (struct run*)v;
3177     r->next = kmem.freelist;
3178     kmem.freelist = r;
3179     if(kmem.use_lock)
3180         release(&kmem.lock);
3181 }
```

Allocating memory to user processes

- System calls like fork, exec allocate memory from OS via kalloc()
- How is address space of process constructed in fork/exec?
 - Start with one page for the outer page directory of child
 - Allocate inner page tables as needed (if contains valid entries)
 - Add page table mappings for kernel code/data (starting at 2GB)
 - Allocate physical frames to store memory contents of process (code/data from executable, empty stack, ..) and map these into page table
- How is address space of process expanded in sbrk?
 - Allocate physical frames for new virtual addresses
 - Add mappings for newly allocated pages in page table

Functions to build page table (1)

- Every page table begins with setting up kernel mappings in setupkvm()
- Outer pgdir allocated
- Kernel mappings defined in “kmap” added to page table by calling “mappages”
- After setupkvm(), user page table mappings added

```
1802 // This table defines the kernel's mappings, which are present in
1803 // every process's page table.
1804 static struct kmap {
1805     void *virt;
1806     uint phys_start;
1807     uint phys_end;
1808     int perm;
1809 } kmap[] = {
1810     { (void*)KERNBASE, 0,          EXTMEM,    PTE_W}, // I/O space
1811     { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},   // kern text+rodata
1812     { (void*)data,      V2P(data),    PHYSTOP,   PTE_W}, // kern data+memory
1813     { (void*)DEVSPACE, DEVSPACE,      0,        PTE_W}, // more devices
1814 };
1815
1816 // Set up kernel part of a page table.
1817 pde_t*
1818 setupkvm(void)
1819 {
1820     pde_t *pgdir;
1821     struct kmap *k;
1822
1823     if((pgdir = (pde_t*)kalloc()) == 0)
1824         return 0;
1825     memset(pgdir, 0, PGSIZE);
1826     if (P2V(PHYSTOP) > (void*)DEVSPACE)
1827         panic("PHYSTOP too high");
1828     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1829         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1830                     (uint)k->phys_start, k->perm) < 0) {
1831             freevm(pgdir);
1832             return 0;
1833         }
1834     return pgdir;
1835 }
```

Functions to build page table (2)

- Page table entries added by “mappages”
 - Arguments: page directory, range of virtual addresses, physical addresses to map to, permissions of the pages
 - For each page, walks page table, get pointer to PTE via function “walkpgdir”, fills it with physical address and permissions

```
1756 // Create PTEs for virtual addresses starting at va that refer to
1757 // physical addresses starting at pa. va and size might not
1758 // be page-aligned.
1759 static int
1760 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1761 {
1762     char *a, *last;
1763     pte_t *pte;
1764
1765     a = (char*)PGROUNDDOWN((uint)va);
1766     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1767     for(;;){
1768         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1769             return -1;
1770         if(*pte & PTE_P)
1771             panic("remap");
1772         *pte = pa | perm | PTE_P;
1773         if(a == last)
1774             break;
1775         a += PGSIZE;
1776         pa += PGSIZE;
1777     }
1778     return 0;
1779 }
```


Functions to build page table (3)

- Function “walkpgdir” walks page table, returns PTE of a virtual address
- Can allocate inner page table if it doesn't exist (depending on value of last arg)

```
1731 // Return the address of the PTE in page table pgdir
1732 // that corresponds to virtual address va.  If alloc!=0,
1733 // create any required page table pages.
1734 static pte_t *
1735 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1736 {
1737     pde_t *pde;
1738     pte_t *pgtab;
1739
1740     pde = &pgdir[PDX(va)];
1741     if(*pde & PTE_P){
1742         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
1743     } else {
1744         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1745             return 0;
1746         // Make sure all those PTE_P bits are zero.
1747         memset(pgtab, 0, PGSIZE);
1748         // The permissions here are overly generous, but they can
1749         // be further restricted by the permissions in the page table
1750         // entries, if necessary.
1751         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
1752     }
1753     return &pgtab[PTX(va)];
1754 }
```

Recap: fork system call implementation

- Parent allocates new process in ptable, copies parent state to child
- Child process set to runnable, scheduler runs it at a later time
- Return value in parent is PID of child, return value in child is set to 0

```
2579 int
2580 fork(void)
2581 {
2582     int i, pid;
2583     struct proc *np;
2584     struct proc *curproc = myproc();
2585
2586     // Allocate process.
2587     if((np = allocproc()) == 0){
2588         return -1;
2589     }
2590
2591     // Copy process state from proc.
2592     if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0){
2593         kfree(np->kstack);
2594         np->kstack = 0;
2595         np->state = UNUSED;
2596         return -1;
2597     }
2598     np->sz = curproc->sz;
2599     np->parent = curproc;
```

```
2600     *np->tf = *curproc->tf;
2601
2602     // Clear %eax so that fork returns 0 in the child.
2603     np->tf->eax = 0;
2604
2605     for(i = 0; i < NOFILE; i++)
2606         if(curproc->ofile[i])
2607             np->ofile[i] = filedup(curproc->ofile[i]);
2608     np->cwd = idup(curproc->cwd);
2609
2610     safestrcpy(np->name, curproc->name, sizeof(curproc->name));
2611
2612     pid = np->pid;
2613
2614     acquire(&ptable.lock);
2615
2616     np->state = RUNNABLE;
2617
2618     release(&ptable.lock);
2619
2620     return pid;
2621 }
```

Fork: copying memory image

- Function “copyvm” called by parent to copy parent memory image to child
 - Create new page table for child
 - Walk through parent memory image page by page and copy it to child
- For each page in parent
 - Fetch PTE, get physical address, permissions
 - Allocate new frame for child, copy contents of parent’s page to new page of child
 - Add a PTE from virtual address to physical address of new page in child page table

```
2032 // Given a parent process's page table, create a copy
2033 // of it for a child.
2034 pde_t*
2035 copyvm(pde_t *pgdir, uint sz)
2036 {
2037     pde_t *d;
2038     pte_t *pte;
2039     uint pa, i, flags;
2040     char *mem;
2041
2042     if((d = setupkvm()) == 0)
2043         return 0;
2044     for(i = 0; i < sz; i += PGSIZE){
2045         if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
2046             panic("copyvm: pte should exist");
2047         if(!(*pte & PTE_P))
2048             panic("copyvm: page not present");
2049         pa = PTE_ADDR(*pte);
2050         flags = PTE_FLAGS(*pte);
2051         if((mem = kalloc()) == 0)
2052             goto bad;
2053         memmove(mem, (char*)P2V(pa), PGSIZE);
2054         if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0) {
2055             kfree(mem);
2056             goto bad;
2057         }
2058     }
2059     return d;
2060
2061 bad:
2062     freevm(d);
2063     return 0;
2064 }
```

Copy-on-write fork

- Real operating systems do copy-on-write: child page table also points to parent pages until either of them modifies it
 - Here, xv6 creates separate memory images for parent and child right away
- Copy-on-write fork (not present in xv6, but easy to do):
 - During fork, new page table allocated to child
 - Child page table entries have physical frame numbers of parent memory image pages only, no copy created for child
 - Parent's memory image is marked as read only
 - When parent or child tries to modify, MMU traps to OS
 - As part of trap handling, separate copy of memory image created
 - Finally, two separate copies of memory image for parent and child

Growing memory image: sbrk

- Initially heap is empty, program “break” is at end of stack
 - sbrk() system call invoked by malloc to expand heap
 - Calls “growproc” to grow memory
- To grow memory, allocuvm allocates new pages, adds mappings into page table for new pages
- Whenever page table updated, must update cr3 register and TLB (done even during context switching)

```
2557 int
2558 growproc(int n)
2559 {
2560     uint sz;
2561     struct proc *curproc = myproc();
2562
2563     sz = curproc->sz;
2564     if(n > 0){
2565         if((sz = allocuvm(curproc->pgdir, sz, sz + n)) == 0)
2566             return -1;
2567     } else if(n < 0){
2568         if((sz = deallocuvm(curproc->pgdir, sz, sz + n)) == 0)
2569             return -1;
2570     }
2571     curproc->sz = sz;
2572     switchuvm(curproc);
2573     return 0;
2574 }
```

allocvm: grow address space


- Walk through new virtual addresses, page by page
- Allocate new frame, add mapping to page table with suitable user permissions
- Similarly deallocvm shrinks memory image, frees up pages

```
1926 int
1927 allocvm(pde_t *pgdir, uint oldsz, uint newsz)
1928 {
1929     char *mem;
1930     uint a;
1931
1932     if(newsz >= KERNBASE)
1933         return 0;
1934     if(newsz < oldsz)
1935         return oldsz;
1936
1937     a = PGROUNDUP(oldsz);
1938     for(; a < newsz; a += PGSIZE){
1939         mem = kalloc();
1940         if(mem == 0){
1941             cprintf("allocvm out of memory\n");
1942             deallocvm(pgdir, newsz, oldsz);
1943             return 0;
1944         }
1945         memset(mem, 0, PGSIZE);
1946         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
1947             cprintf("allocvm out of memory (2)\n");
1948             deallocvm(pgdir, newsz, oldsz);
1949             kfree(mem);
1950             return 0;
1951         }
1952     }
1953     return newsz;
1954 }
```

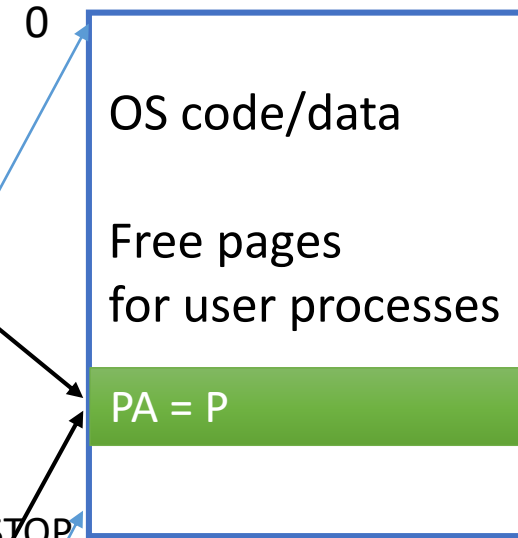
Maximum addressable memory in xv6

- PA=P is initially mapped into kernel address space at VA=P+2GB
- When assigned to user, P is assigned another VA=U (<2GB)
- Kernel and user access same memory using different virtual addresses
- Every physical address may be mapped to 2 virtual addresses in xv6
- Max virtual address is 4GB, so xv6 can only handle max physical address 2GB
- Real kernels deal with this better, e.g., remove kernel VA-PA mapping once assigned to user process

Virtual address space



Physical address space



VA = U

KERNBASE

VA = P + 2GB

KERNBASE +
PHYSTOP

PHYSTOP

Exec system call (1)

- Read ELF binary file from disk into memory
- Start with new page table (not overwriting old page table)
- Use function “loadvm” to load executable from disk to memory

```
6609 int
6610 exec(char *path, char **argv)
6611 {
6612     char *s, *last;
6613     int i, off;
6614     uint argc, sz, sp, ustack[3+MAXARG+1];
6615     struct elfhdr elf;
6616     struct inode *ip;
6617     struct proghdr ph;
6618     pde_t *pgdir, *oldpgdir;
6619     struct proc *curproc = myproc();
6620
6621     begin_op();
6622
6623     if((ip = namei(path)) == 0){
6624         end_op();
6625         cprintf("exec: fail\n");
6626         return -1;
6627     }
6628     ilock(ip);
6629     pgdir = 0;
6630
6631     // Check ELF header
6632     if(readi(ip, (char*)&elf, 0, sizeof(elf)) != sizeof(elf))
6633         goto bad;
6634     if(elf.magic != ELF_MAGIC)
6635         goto bad;
6636
6637     if((pgdir = setupkvm()) == 0)
6638         goto bad;
```

```
6640     // Load program into memory.
6641     sz = 0;
6642     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
6643         if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
6644             goto bad;
6645         if(ph.type != ELF_PROG_LOAD)
6646             continue;
6647         if(ph.memsz < ph.filesz)
6648             goto bad;
6649         if(ph.vaddr + ph.memsz < ph.vaddr)
6650             goto bad;
6651         if((sz = allocvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
6652             goto bad;
6653         if(ph.vaddr % PGSIZE != 0)
6654             goto bad;
6655         if(loadvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
6656             goto bad;
6657     }
6658     iunlockput(ip);
6659     end_op();
6660     ip = 0;
```


Exec system call (2)

- Function allocvm allocates new memory frame, updates page table entries
- Function loadvm reads the corresponding part of executable from disk into the allocated memory frame
- Calls to allocvm and loadvm repeated for each segment of executable

```
1900 // Load a program segment into pgdir.  addr must be page-aligned
1901 // and the pages from addr to addr+sz must already be mapped.
1902 int
1903 loadvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint sz)
1904 {
1905     uint i, pa, n;
1906     pte_t *pte;
1907
1908     if((uint) addr % PGSIZE != 0)
1909         panic("loadvm: addr must be page aligned");
1910     for(i = 0; i < sz; i += PGSIZE){
1911         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
1912             panic("loadvm: address should exist");
1913         pa = PTE_ADDR(*pte);
1914         if(sz - i < PGSIZE)
1915             n = sz - i;
1916         else
1917             n = PGSIZE;
1918         if(readi(ip, P2V(pa), offset+i, n) != n)
1919             return -1;
1920     }
1921     return 0;
1922 }
```

Exec system call (3)

- After executable is copied to memory image, allocate 2 pages for stack (one is guard page, permissions cleared, access will trap)
- Push exec arguments onto user stack for main function of new program

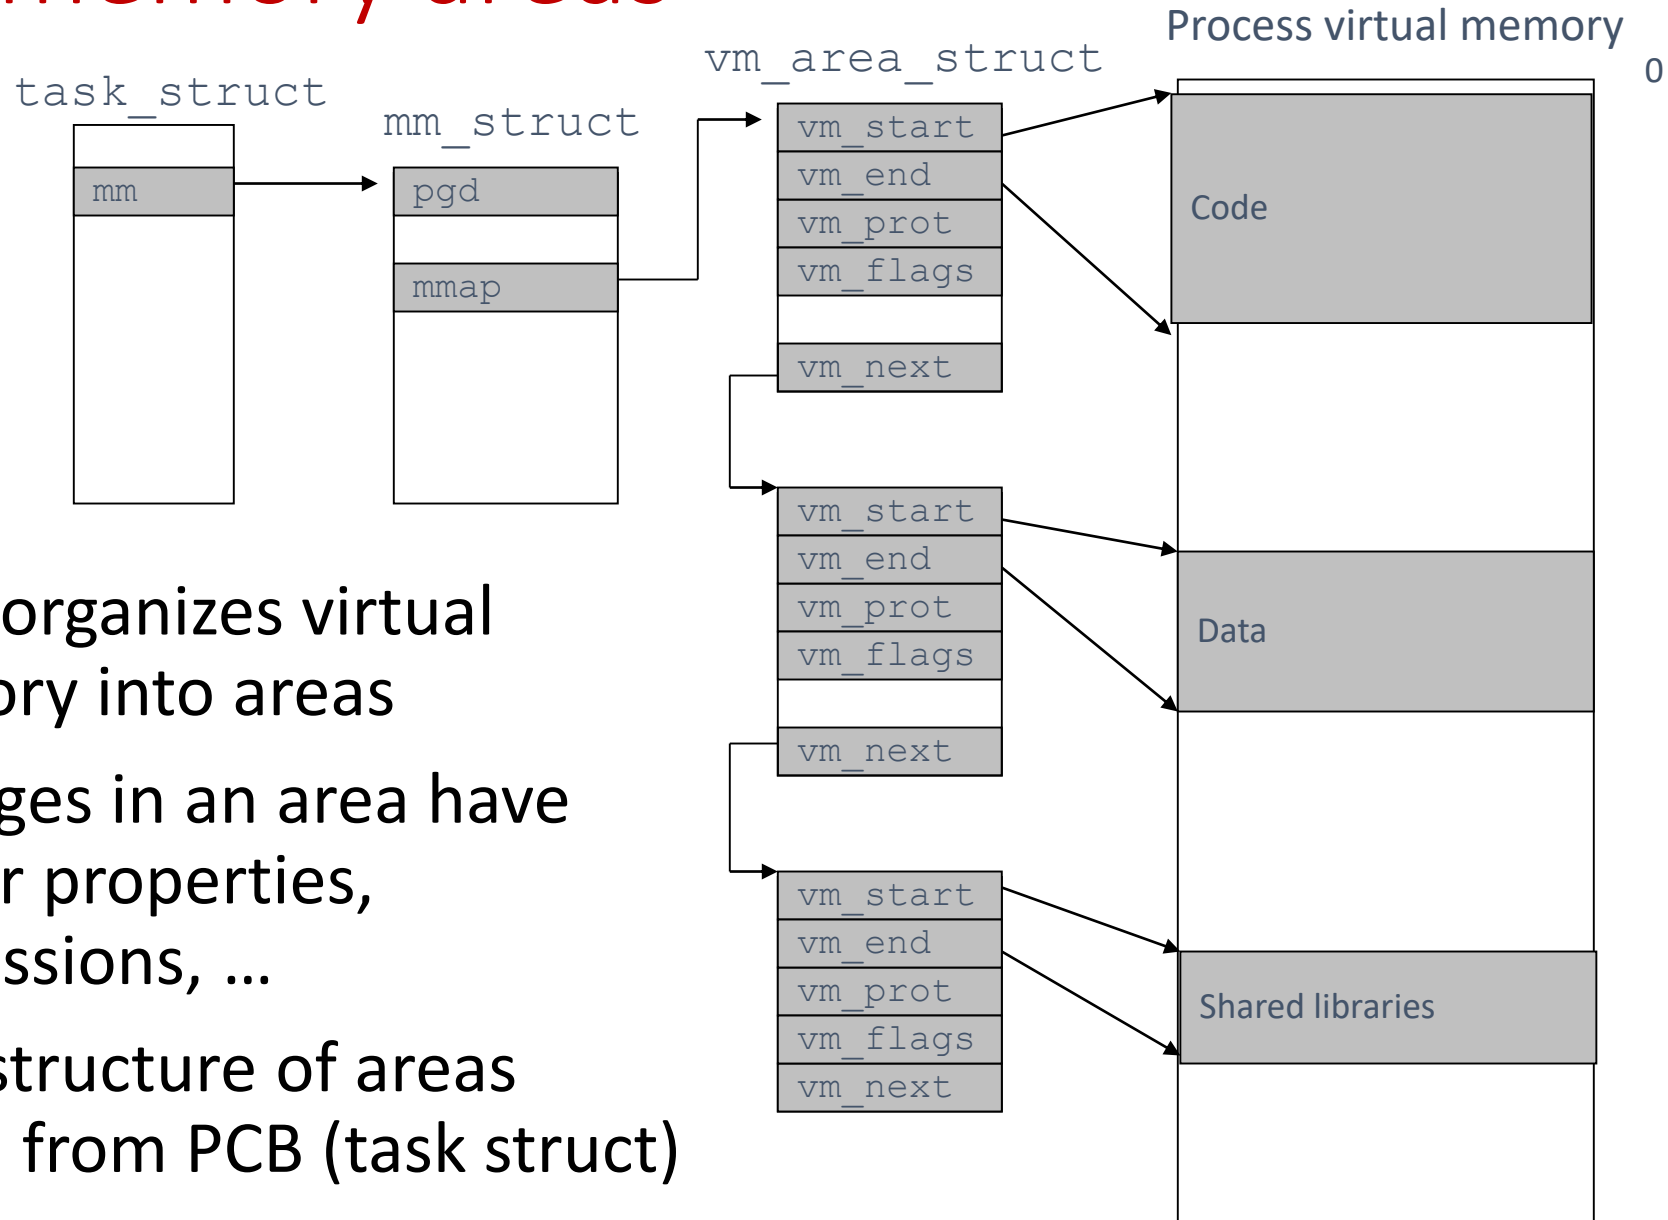
```
6662 // Allocate two pages at the next page boundary.
6663 // Make the first inaccessible. Use the second as the user stack.
6664 sz = PGROUNDUP(sz);
6665 if((sz = allocuvm(pgdir, sz, sz + 2*PGSIZE)) == 0)
6666     goto bad;
6667 clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
6668 sp = sz;
6669
6670 // Push argument strings, prepare rest of stack in ustack.
6671 for(argc = 0; argv[argc]; argc++) {
6672     if(argc >= MAXARG)
6673         goto bad;
6674     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
6675     if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
6676         goto bad;
6677     ustack[3+argc] = sp;
6678 }
6679 ustack[3+argc] = 0;
6680
6681 ustack[0] = 0xffffffff; // fake return PC
6682 ustack[1] = argc;
6683 ustack[2] = sp - (argc+1)*4; // argv pointer
6684
6685 sp -= (3+argc+1) * 4;
6686 if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
6687     goto bad;
6688
```

Exec system call (4)

- If no errors so far, switch to new page table that is pointing to new memory image
 - If any error, go back to old memory image (exec returns with error)
- Set eip in trapframe to start at entry point of new program
 - Returning from trap, process will run new executable

```
6689 // Save program name for debugging.
6690 for(last=s=path; *s; s++)
6691     if(*s == '/')
6692         last = s+1;
6693 safestrcpy(curproc->name, last, sizeof(curproc->name));
6694
6695 // Commit to the user image.
6696 oldpgdir = curproc->pgdir;
6697 curproc->pgdir = pgdir;
6698 curproc->sz = sz;
6699 curproc->tf->eip = elf.entry; // main
6700 curproc->tf->esp = sp;
6701 switchvm(curproc);
6702 freevm(oldpgdir);
6703 return 0;
6704
6705 bad:
6706 if(pgdir)
6707     freevm(pgdir);
6708 if(ip){
6709     iunlockput(ip);
6710     end_op();
6711 }
6712 return -1;
6713 }
```

Linux memory areas



- Linux organizes virtual memory into areas
- All pages in an area have similar properties, permissions, ...
- Data structure of areas linked from PCB (task struct)